

main



TensorFlow_Tutorials / TensorFlow_2.x / Chapter4_LeNet.ipynb



hy23 Folder restructuring

History

0 contributors

10.2 MB



```
In [1]: import tensorflow as tf
        from tensorflow.keras import datasets, layers, models, optimizers
```

```
In [2]: # Load the TensorBoard notebook extension
        # https://www.tensorflow.org/tensorboard/get_started

        %load_ext tensorboard

        import datetime
```

```
In [3]: # Clear any logs from previous runs
        !rm -rf /content/logs
```

```
In [4]: # network and training
        EPOCHS           = 10
        BATCH_SIZE       = 128
        VERBOSE          = 0
        OPTIMIZER         = tf.keras.optimizers.Adam()
        VALIDATION_SPLIT = 0.95

        # input image dimensions
        IMG_ROWS, IMG_COLS = 28, 28
        INPUT_SHAPE        = (IMG_ROWS, IMG_COLS, 1)

        # number of outputs = number of digits
        NB_CLASSES         = 10
```

```
In [5]: def build(input_shape, classes):
        model = models.Sequential()
        model.add(layers.Conv2D(filters=20, kernel_size=(5,5), padding='same',
                                activation='relu', input_shape=input_shape))
        model.add(layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))
        model.add(layers.Conv2D(filters=50, kernel_size=(5,5), padding='same',
                                activation='relu'))
        model.add(layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))
        model.add(layers.Flatten())
        model.add(layers.Dense(500, activation='relu'))
        # a softmax classifier
        model.add(layers.Dense(classes, activation='softmax'))
        return model
```

```
In [6]: # Load data
        (x_train, y_train), (x_test, y_test) = datasets.mnist.load_data()

        print("Original shape of x_train is {}".format(x_train.shape))
```

Original shape of x_train is (60000, 28, 28)

```
In [7]: # reshape
        x_train = x_train.reshape((60000, 28, 28, 1))
        x_test  = x_test.reshape((10000, 28, 28, 1))

        # normalize
        x_train, x_test = x_train / 255.0, x_test / 255.0

        # cast
        x_train = x_train.astype('float32')
        x_test  = x_test.astype('float32')
```

```
In [8]: print("Original shape of y_train is {}".format(y_train.shape))
```

Original shape of y_train is (60000,)

```
In [9]: y_train = tf.keras.utils.to_categorical(y_train, NB_CLASSES)
        y_test  = tf.keras.utils.to_categorical(y_test, NB_CLASSES)

        print("Shape of y_train is {}".format(y_train.shape))
```

Shape of y_train is (60000, 10)

```
In [10]: in_shape = x_train.shape[1:]

# get the model
model = build(in_shape, NB_CLASSES)
```

```
In [11]: model.compile(optimizer=OPTIMIZER, loss='categorical_crossentropy',
                    metrics=['accuracy'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 20)	520
max_pooling2d (MaxPooling2D)	(None, 14, 14, 20)	0
conv2d_1 (Conv2D)	(None, 14, 14, 50)	25050
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 50)	0
flatten (Flatten)	(None, 2450)	0
dense (Dense)	(None, 500)	1225500
dense_1 (Dense)	(None, 10)	5010
Total params: 1,256,080		
Trainable params: 1,256,080		
Non-trainable params: 0		

```
In [12]: # use TensorBoard
callbacks = [
    # write TensorBoard logs to './logs' directory
    tf.keras.callbacks.TensorBoard(log_dir='./logs')
]
```

```
In [13]: # train the model
history = model.fit(x=x_train, y=y_train, batch_size=BATCH_SIZE, epochs=EPOCHS,
                    verbose=VERBOSE, callbacks=callbacks,
                    validation_split=VALIDATION_SPLIT)
```

```
In [14]: score = model.evaluate(x_test, y_test, verbose=VERBOSE)
```

```
In [15]: print("test loss is {}".format(score[0]))
print("test accu is {}".format(score[1]))
```

test loss is 0.09741149842739105
test accu is 0.9700000286102295

TensorBoard

```
In [16]: # How to use TensorBoard?
# https://www.tensorflow.org/tensorboard/get_started
%tensorboard --logdir logs/train
```

Reusing TensorBoard on port 6007 (pid 713), started 4:00:51 ago. (Use '!kill 713' to kill it.)

```
In [17]: # How to visualize feature maps.

for layer in model.layers:
    # from model.summary(), we know the names of the 'conv' layers
    if 'conv2d' not in layer.name:
        continue
    # get filter weights
    filters, biases = layer.get_weights()
    print(layer.name, filters.shape)
```

conv2d (5, 5, 1, 20)
conv2d_1 (5, 5, 20, 50)

```
conv2d_1 (5, 5, 20, 20)
```


Visualize Feature Maps

```
In [18]: # https://machinelearningmastery.com/how-to-visualize-filters-and-feature-maps-in-convolutional-neural-n

# The idea of visualizing a feature map for a specific input image would be to
# understand what features of the input are detected or preserved in the feature
# maps. The expectation would be that the feature maps close to the input detect
# small or fine-grained detail, whereas feature maps close to the output of the
# model capture more general features.
```

```
In [19]: from matplotlib import pyplot as plt
```

```
In [20]: plt.imshow(x_test[0].reshape(28,28), cmap='gray')
plt.show()
```



```
In [21]: layer_outputs = [layer.output for layer in model.layers]
layer_outputs
```

```
Out[21]: [<KerasTensor: shape=(None, 28, 28, 20) dtype=float32 (created by layer 'conv2d')>,
<KerasTensor: shape=(None, 14, 14, 20) dtype=float32 (created by layer 'max_pooling2d')>,
<KerasTensor: shape=(None, 14, 14, 50) dtype=float32 (created by layer 'conv2d_1')>,
<KerasTensor: shape=(None, 7, 7, 50) dtype=float32 (created by layer 'max_pooling2d_1')>,
<KerasTensor: shape=(None, 2450) dtype=float32 (created by layer 'flatten')>,
<KerasTensor: shape=(None, 500) dtype=float32 (created by layer 'dense')>,
<KerasTensor: shape=(None, 10) dtype=float32 (created by layer 'dense_1')>]
```

```
In [22]: feature_map_model = tf.keras.models.Model(inputs=model.input, outputs=layer_outputs)
```

```
In [23]: sample = x_test[0]
print("shape of the sample is {}".format(sample.shape))
sample = sample.reshape(1, 28, 28, 1)
print("shape of the sample is {}".format(sample.shape))
```

```
shape of the sample is (28, 28, 1)
shape of the sample is (1, 28, 28, 1)
```

```
In [24]: fmap = feature_map_model.predict(sample)
```

```
In [25]: print("type of the fmap is {}".format(type(fmap)))
print("shape of the fmap is {}".format(len(fmap)))
```

```
type of the fmap is <class 'list'>
shape of the fmap is 7
```


```
In [26]: # each fmap corresponds to each layer.
# we know that layer0 is a convolution layer.

fmap_0 = fmap[0]
print("shape of the feature map 0 is {}".format(fmap_0.shape))
```

```
shape of the feature map 0 is (1, 28, 28, 20)
```

```
In [27]: # example plot
plt.imshow(fmap_0[0, :, :, 1], cmap='gray')
```

Out[27]: <matplotlib.image.AxesImage at 0x7fda77646450>




```
In [28]: # Let us plot all the 20 feature maps of convolution layer 0
no_of_plots_in_x = 5
no_of_plots_in_y = 4

index = 1
for x in range(no_of_plots_in_x):
    for y in range(no_of_plots_in_y):
        subplot = plt.subplot(no_of_plots_in_x, no_of_plots_in_y, index)
        subplot.set_xticks([])
        subplot.set_yticks([])
        # buf = sprintf("feature map # %d", index)
        # buf = f"depth #{index-1:d}"
        # subplot.set_title(buf)
        plt.imshow(fmap_0[0, :, :, index-1], cmap='gray')
        index += 1

print("Feature maps of convolution layer 0")
plt.show()
```

Feature maps of convolution layer 0



```
In [29]: fmap_2 = fmap[2]
print("shape of the feature map 0 is {}".format(fmap_2.shape))
```

shape of the feature map 0 is (1, 14, 14, 50)


```
In [30]: # Let us plot all the 50 feature maps of convolution layer 0
no_of_plots_in_x = 7
no_of_plots_in_y = 7

index = 1
for x in range(no_of_plots_in_x):
    for y in range(no_of_plots_in_y):
        subplot = plt.subplot(no_of_plots_in_x, no_of_plots_in_y, index)
        subplot.set_xticks([])
        subplot.set_yticks([])
        # buf = sprintf("feature map # %d", index)
        # buf = f"depth #{index-1:d}"
        # subplot.set_title(buf)
        plt.imshow(fmap_2[0, :, :, index-1], cmap='gray')
        index += 1

print("Feature maps of convolution layer 2")
plt.show()
```

Feature maps of convolution layer 2





In [30]:

